




Logistic Regression




- If i want to predict, weather a person having smart phone or not with respect to linear line



- Lets fit a line
ideally the value should onle be zero and one




if we decide a threshold,
if the value of age is below the
perticular pint the person does
have a smart phone




the moment we add some data point,
the slope of the line changes and so
that the threshold

It does not give a right way to classify a problem




also, when we extend the line, it
does not mean anything.....

**if classification, below 0 and above 1
does not have any meaning**



$$y = mx + c$$

here, i actually need a function of y which transform the linear line to cover all the data points



$$Z = \beta X + b$$

$$\hat{Y} = Q(Z)$$

$$Q(Z) = \frac{1}{1 + e^{-Z}}$$

$$\hat{Y} = \frac{1}{1 + e^{-Z}}$$


$$\hat{Y} = \frac{1}{1 + e^{-(\beta X + b)}}$$

The prediction through line we get are continuous in nature, so we can use them as the probabilities to make prediction.....

- if the data point is extreme right than probabilities are close to one and visa-versa


$$\text{Log loss} = \frac{1}{N} \sum_{i=1}^N -(y_i * \log(\hat{Y}_i) + (1 - y_i) * \log(1 - \hat{Y}_i))$$

When $y=1$



$$J = \frac{1}{N} \sum_{i=1}^N -(y_i * \log(\hat{Y}_i) + (1 - y_i) * \log(1 - \hat{Y}_i))$$

When $y=0$



$$J = \frac{1}{N} \sum_{i=1}^N -(y_i * \log(\hat{Y}_i) + (1 - y_i) * \log(1 - \hat{Y}_i))$$

Gradient Descent in Logistic Regression

$$J = \frac{\sum_{i=1}^n - (Y_i \log(\hat{Y}_i) + (1-Y_i) \log(1-\hat{Y}_i))}{n}$$

$$\beta = \beta - \alpha G_\beta$$

$$b = b - \alpha G_b$$

$$G_\beta = \frac{\partial(J)}{\partial \beta} = \frac{-2 \sum_{i=1}^n (\hat{Y}_i - Y_i) X_i}{n} = \frac{-2}{n} \sum_{i=1}^n \left(\frac{1}{1 + e^{-(\beta X_i + b)}} - Y_i \right) X_i$$

$$G_b = \frac{\partial(J)}{\partial b} = \frac{-2 \sum_{i=1}^n (\hat{Y}_i - Y_i)}{n} = \frac{-2}{n} \sum_{i=1}^n \left(\frac{1}{1 + e^{-(\beta X_i + b)}} - Y_i \right)$$


Optimization Algorithms:

Gradient Descent

Conjugate gradient

BFGS

L - BFGS



Advantages:

- No need to manually pick α
- Often faster than gradient descent.

Disadvantages:

- More complex

Multiclass Classification

One Vs. All method : Building a model with one class versus all (Not that class)


class- A Vs. Not class - A
class- B Vs. Not class - B
class- C Vs. Not class - C

Get the
Probability
outcomes

ID	AGE	Class A	ID	AGE	Class B	ID	AGE	Class C
1	2	1	1	2	0	1	2	0
2	33	0	2	33	0	2	33	1
3	7	1	3	7	0	3	7	0
4	25	0	4	25	0	4	25	1
5	4	1	5	4	0	5	4	0
6	48	0	6	48	1	6	48	0


ID	AGE	Number of phones	P(A)	P(B)	P(C)
1	2	A	0.92	0.12	0.09
2	33	C	0.33	0.17	0.94
3	7	A	0.83	0.11	0.10
4	25	C	0.14	0.40	0.88
5	4	A	0.87	0.23	0.21
6	48	B	0.21	0.79	0.39

Which ever class has
higher probability, we
will make that
prediction




Bias: Error in training \longrightarrow leads to **UNDERFITTING**: model has not fitted training data


Variance: Error in \longrightarrow leads to **OVERFITTING**: model has overly trained on training data, not generalized model



$\rightarrow \theta_0 + \theta_1 x$
"Underfit" "High bias"



$\rightarrow \theta_0 + \theta_1 x + \theta_2 x^2$
"Just right"



$\rightarrow \theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$
"Overfit" "High variance"

Underfittin




High bias,
High

Overfitting



Low bias,
High

Bias: Bias is the difference between the average prediction of our model and the correct value which we are trying to predict.

- Model with the high bias pays little attention to the training data and oversimplifies the model.
 - It always leads to high error in training and testing data.
- 

Variance: Variance is the variability of model prediction for a given data point or value which tells us the spread of the data.


- Model with high variance pays a lot of attention to training data and does not generalize on the data which it has not seen before.
- As a result, the model performs very well on training data, but has high error in testing data.

In supervised learning, **underfitting** happens when a model is unable to capture the underlying pattern of the data. These models usually have high bias and low variance. It happens when we have very less amount of data to build an accurate model or when we try to build a linear model with a nonlinear data. Also, these kind of models are very simple to capture the complex patterns in data like Linear and logistic regression.

In supervised learning, **overfitting** happens when our model captures the noise along with the underlying pattern in data. It happens when we train our model a lot over a noisy dataset. These models have low bias and high variance. These models are very complex like Decision trees which are prone to overfitting.

Bias-Variance Tradeoff

$$\text{Total Error} = \text{Bias}^2 + \text{Variance} + \text{Irreducible Error}$$



Addressing OVERFITTING

1) Reduce the number of features:


- Manually select which features to keep.
- Use a model selection algorithm (studied later in the course).

2) Regularization

- Keep all the features, but reduce the magnitude of parameters θ_j .
- Regularization works well when we have a lot of slightly useful features.


Regularization

Intuition



Size of house

$$\theta_0 + \theta_1 x + \theta_2 x^2$$



Size of house

$$\theta_0 + \theta_1 x + \theta_2 x^2 + \theta_3 x^3 + \theta_4 x^4$$

To overcome the **OVERFITTING** we can convert this polynomial equation to quadratic equation (red line)

Suppose we penalize and make θ_3, θ_4 really small.

$$\min_{\theta} \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + 1000 \theta_3^2 + 1000 \theta_4^2$$

$$\theta_3 \approx 0$$

$$\theta_4 \approx 0$$

Regularization.

Small values for parameters $\theta_0, \theta_1, \dots, \theta_n$

- “Simpler” hypothesis ←
- Less prone to overfitting

$$J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

$\theta_1, \theta_2, \theta_3, \dots, \theta_{100}$

We have to add regularization term to the cost function

$$\min_{\theta} J(\theta) = \frac{1}{2m} \left[\sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2 + \lambda \sum_{j=1}^n \theta_j^2 \right]$$

Regularization Parameter;


Lambda choose the tradoff between two different goals

i) We would like to fit training data well

ii) We would like to keep parameters small

The λ , or lambda, is the **regularization parameter**. It determines how much the costs of our theta parameters are inflated.

What if λ is set to an extremely large value (perhaps far too large for our problem, say $\lambda = 10^{10}$)?



$$\theta_1, \theta_2, \theta_3, \theta_4$$

$$\theta_1 \approx 0, \theta_2 \approx 0$$

$$\theta_3 \approx 0, \theta_4 \approx 0$$

$$h_{\theta}(x) = \theta_0$$

If lambda is chosen to be too large, it may smooth out the function too much and cause underfitting.

Regularization in Linear Regression

Repeat {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} \quad (j = \text{red X}, 1, 2, 3, \dots, n) \}$$

$$\theta_j := \theta_j - \alpha \left[\frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} + \frac{\lambda}{m} \theta_j \right]$$

(j = ~~red X~~, 1, 2, 3, ..., n)

$\frac{\partial}{\partial \theta_j} J(\theta)$ regularized

$$\theta_j := \theta_j (1 - \alpha \frac{\lambda}{m}) - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

$$(1 - \alpha \frac{\lambda}{m}) < 1$$

The first term in the above equation, $1 - \alpha \frac{\lambda}{m}$ will always be less than 1. Intuitively you can see it as reducing the value of θ_j by some amount on every update. Notice that the second term is now exactly the same as it was before.